

In [1]:

```
1 import pandas as pd
2 from random import gauss as gs
3 import pandas as pd
4 import numpy as np
5 import matplotlib.pyplot as plt
6 from sklearn.metrics import mean_squared_error
7 from sklearn.linear_model import LinearRegression
8 from sklearn.model_selection import TimeSeriesSplit
9 import matplotlib.dates as mdates
10 import matplotlib as mpl
11 import seaborn as sns
12 from math import sqrt
13
14 import itertools
15 #from pmdarima import auto_arima
16
17 #statsmodels
18
19 from statsmodels.tsa.arima.model import ARIMA
20 from statsmodels.tsa.stattools import acf, pacf, adfuller
21 from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
22 from statsmodels.tsa.statespace.sarimax import SARIMAX
23 from statsmodels.graphics.tsaplots import plot_predict
24 import statsmodels.api as sm
25
26 %matplotlib inline
```

In [2]:

```

1 crude_oil = pd.read_csv("data/crude_oil.csv")
2 gold = pd.read_csv("data/gold.csv")
3 dow_jones = pd.read_csv("data/dow_jones.csv")
4 fed_funds = pd.read_csv("data/fed_funds.csv")
5
6 crude_oil['Crude Oil'] = crude_oil['real']
7 crude_oil = crude_oil.drop(['real', 'nominal'], axis = 1)
8 crude_oil['date'] = pd.to_datetime(crude_oil['date'])
9 crude_oil.set_index('date', inplace = True)
10
11
12 gold['Gold'] = gold['real']
13 gold = gold.drop(['real', 'nominal'], axis = 1)
14 gold['date'] = pd.to_datetime(gold['date'])
15 gold.set_index('date', inplace = True)
16
17 dow_jones['Dow Jones'] = dow_jones['real']
18 dow_jones = dow_jones.drop(['real', 'nominal'], axis = 1)
19 dow_jones['date'] = pd.to_datetime(dow_jones['date'])
20 dow_jones.set_index('date', inplace = True)
21
22 fed_funds['date'] = fed_funds['DATE']
23 fed_funds['Fed Funds'] = fed_funds['FEDFUNDS']
24 fed_funds = fed_funds.drop(['DATE', 'FEDFUNDS'], axis = 1)
25 fed_funds['date'] = pd.to_datetime(fed_funds['date'])
26 fed_funds.set_index('date', inplace = True)
27
28 crude_oil

```

Out[2]:

Crude Oil	
date	
1946-01-01	18.79
1946-02-01	18.89
1946-03-01	18.69
1946-04-01	20.18
1946-05-01	20.07
...	...
2022-03-01	101.98
2022-04-01	105.84
2022-05-01	114.67
2022-06-01	159.57
2022-07-01	123.71

919 rows × 1 columns

In [3]:

```
1 occidental = pd.read_csv("data/OXY Historical Data.csv")
2 occidental['date'] = pd.to_datetime(occidental['Date'], format = "%b %y")
3 occidental['OXY Price'] = occidental['Price']
4 occidental = occidental.drop(['Date', 'Open', 'High', 'Low', 'Vol.', 'Change %', 'Price'])
5 occidental = occidental.set_index('date')
6 occidental
```

Out[3]:

OXY Price	
date	
2022-07-01	59.52
2022-06-01	58.88
2022-05-01	69.31
2022-04-01	55.09
2022-03-01	56.74
...	...
1980-08-01	13.19
1980-07-01	12.59
1980-06-01	12.83
1980-05-01	12.71
1980-04-01	11.33

508 rows × 1 columns

In [4]:

```

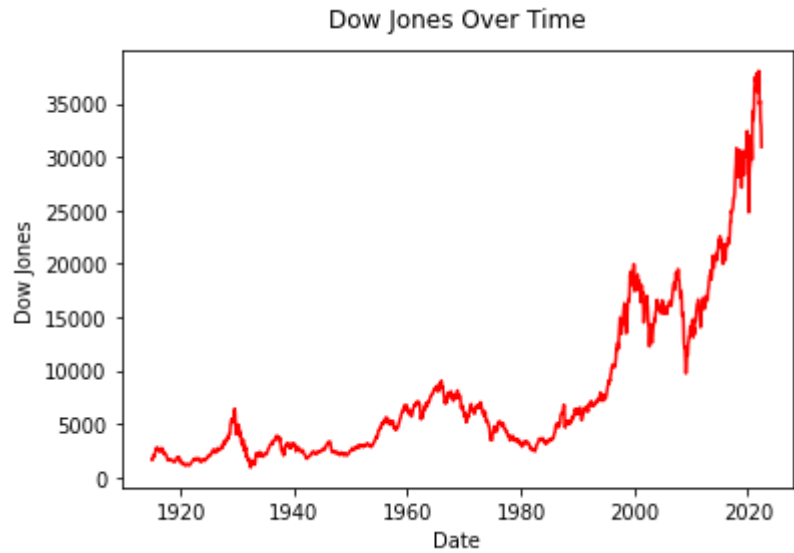

1 all_data = dow_jones.copy()
2 all_data['Gold'] = gold['Gold']
3 #all_data = all_data[400:]
4 #all_data['Crude Oil'] = crude_oil['Crude Oil']
5
6 all_data = pd.concat([all_data, crude_oil, fed_funds, occidental], axis=1)
7 #all_data = pd.concat([all_data, fed_funds], axis=1)
8 #all_data['Fed Funds'] = fed_funds['Fed Funds']
9
10
11 #all_data.set_index('date', inplace = True)
12
13
14 fig, ax1 = plt.subplots()
15 ax1.set_title('OXY Price and Crude Oil Price Over Time', pad = 12, fontsize = 15)
16 color = 'blue'
17 ax1.set_xlabel('Date')
18 ax1.set_ylabel('OXY Price', color = color)
19 ax1.plot(all_data.index, all_data['OXY Price'], color = color)
20 ax1.tick_params(axis = 'y', labelcolor = color)
21 ax1.yaxis.set_major_formatter(mpl.ticker.StrMethodFormatter('${x:,.0f}'))
22
23 # Adding Twin Axes to plot using dataset_2
24 ax2 = ax1.twinx()
25
26 color = 'black'
27 ax2.set_ylabel('Crude Oil', color = color)
28 ax2.plot(all_data.index, all_data['Crude Oil'], color = color)
29 ax2.tick_params(axis = 'y', labelcolor = color)
30 ax2.yaxis.set_major_formatter(mpl.ticker.StrMethodFormatter('${x:,.0f}'))
31
32
33 fig, ax3 = plt.subplots()
34 ax3.set_title('Dow Jones Over Time', pad = 12)
35 ax3.set_xlabel('Date')
36 color = 'red'
37 ax3.set_ylabel('Dow Jones')
38 ax3.plot(all_data.index, all_data['Dow Jones'], color = color)
39
40
41 fig, ax4 = plt.subplots()
42 ax4.set_title('Fed Funds Rate Over Time', pad = 12)
43 ax4.set_xlabel('Date')
44 ax4.set_ylabel('Fed')
45 ax4.plot(all_data.index, all_data['Fed Funds'])
46 ax4.yaxis.set_major_formatter(mpl.ticker.StrMethodFormatter('{x:,.0f}%'))
47
48 fig, ax5 = plt.subplots()
49 ax5.set_title('Crude Oil Price Over Time', pad = 12, fontsize = 15)
50 color = 'black'
51 ax5.set_ylabel('Crude Oil', color = color)
52 ax5.plot(all_data.index, all_data['Crude Oil'], color = color)
53 ax5.tick_params(axis = 'y', labelcolor = color)
54 ax5.yaxis.set_major_formatter(mpl.ticker.StrMethodFormatter('${x:,.0f}'))
55 all_data
56
57

```


Out[4]:

	Dow Jones	Gold	Crude Oil	Fed Funds	OXY Price
date					
1915-01-01	1636.27	557.10	NaN	NaN	NaN
1915-02-01	1608.23	562.68	NaN	NaN	NaN
1915-03-01	1796.01	568.36	NaN	NaN	NaN
1915-04-01	2098.13	562.68	NaN	NaN	NaN
1915-05-01	1881.39	557.10	NaN	NaN	NaN
...	...	...	...	...	...
2022-03-01	35267.88	1986.24	101.98	0.20	56.74
2022-04-01	33339.96	1932.73	105.84	0.33	55.09
2022-05-01	32990.12	1847.26	114.67	0.77	69.31
2022-06-01	30946.99	1824.80	159.57	NaN	58.88
2022-07-01	NaN	NaN	123.71	NaN	59.52


1291 rows × 5 columns



Fed Funds Rate Over Time



Crude Oil Price Over Time



In [5]:

```

1 with sns.axes_style('darkgrid'):
2
3     f, ax = plt.subplots(figsize=(9,9))
4
5     mask = np.triu(np.ones_like(all_data.corr(), dtype=np.bool))
6
7     plt.xticks(fontsize = 15)
8
9     plt.yticks(fontsize = 15)
10
11    sns.set(font_scale=1.4)
12
13    heatmap = sns.heatmap(all_data.corr(), annot = True, mask = mask, cmap = "BrBG")
14
15    heatmap.set_title("Correlation Heatmap of Economic Factors", fontdict={'fontsize':

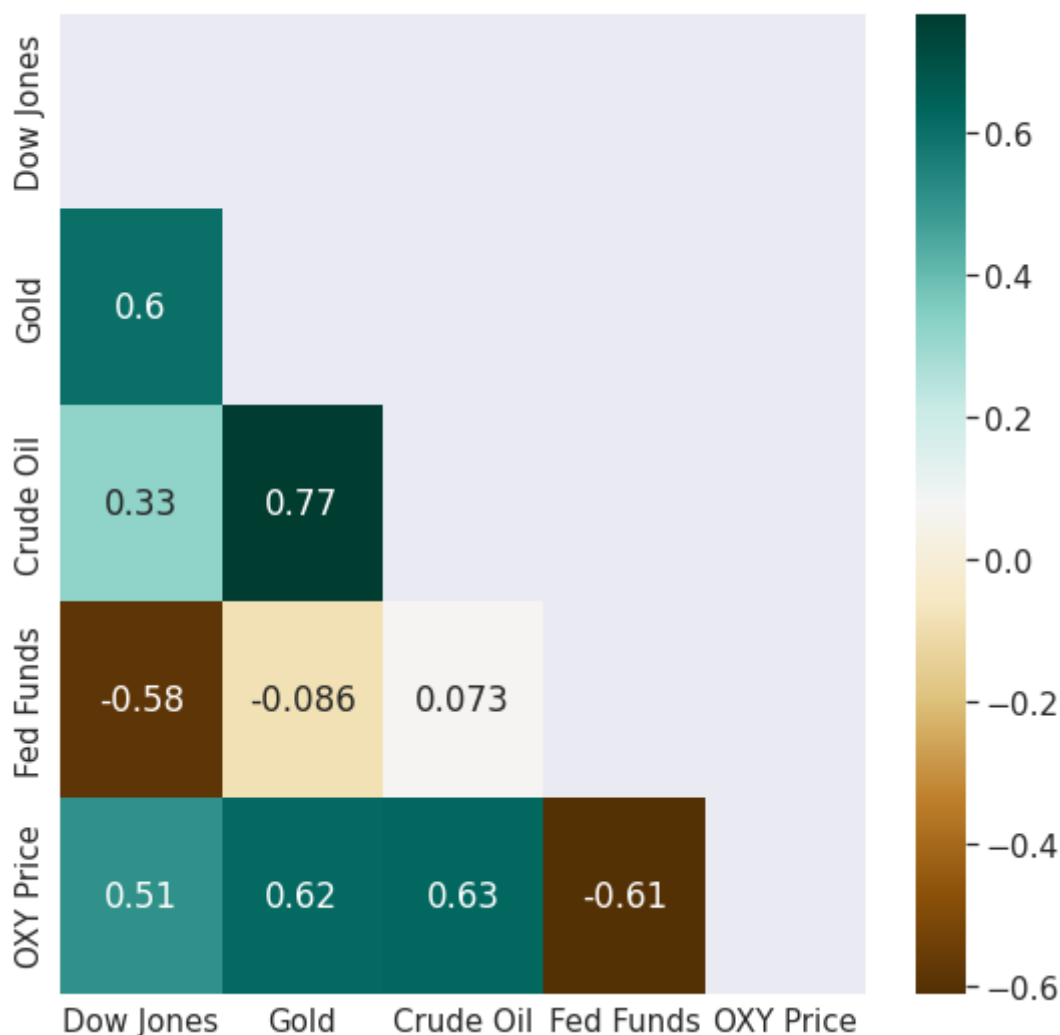
```

C:\Users\wjsdn\AppData\Local\Temp\ipykernel\_31800\504848525.py:5: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool\_` here.

Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations> (<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>)

```
mask = np.triu(np.ones_like(all_data.corr(), dtype=np.bool))
```

## Correlation Heatmap of Economic Factors



In [6]:

```
1 all_data['Crude Oil']
```

Out[6]:

date	
1915-01-01	NaN
1915-02-01	NaN
1915-03-01	NaN
1915-04-01	NaN
1915-05-01	NaN

...

2022-03-01	101.98
2022-04-01	105.84
2022-05-01	114.67
2022-06-01	159.57
2022-07-01	123.71

Freq: MS, Name: Crude Oil, Length: 1291, dtype: float64



In [7]:

```

1 train_df = all_data['1990-07-01':'2021-05-01']
2 test_df = all_data['2021-05-01':'2022-07-01']
3
4 test_df

```

Out[7]:

	Dow Jones	Gold	Crude Oil	Fed Funds	OXY Price
date					
2021-05-01	37498.98	2068.55	72.02	0.06	25.96
2021-06-01	37124.70	1906.24	79.05	0.08	31.27
2021-07-01	37415.89	1942.92	79.20	0.10	26.10
2021-08-01	37765.26	1939.27	73.16	0.09	25.69
2021-09-01	36077.62	1872.60	79.98	0.08	29.58
2021-10-01	37861.27	1885.58	88.33	0.08	33.53
2021-11-01	36276.87	1868.27	69.62	0.08	29.65
2021-12-01	38082.54	1916.37	78.82	0.08	28.99
2022-01-01	36537.13	1867.96	91.68	0.08	37.67
2022-02-01	34909.38	1957.72	98.59	0.08	43.73
2022-03-01	35267.88	1986.24	101.98	0.20	56.74
2022-04-01	33339.96	1932.73	105.84	0.33	55.09
2022-05-01	32990.12	1847.26	114.67	0.77	69.31
2022-06-01	30946.99	1824.80	159.57	NaN	58.88
2022-07-01	NaN	NaN	123.71	NaN	59.52

In [8]:

```
1 test_df_oil = test_df['Crude Oil']  
2 test_df_oil = pd.DataFrame(test_df_oil)  
3 test_df_oil
```

Out[8]:

Crude Oil	
date	
2021-05-01	72.02
2021-06-01	79.05
2021-07-01	79.20
2021-08-01	73.16
2021-09-01	79.98
2021-10-01	88.33
2021-11-01	69.62
2021-12-01	78.82
2022-01-01	91.68
2022-02-01	98.59
2022-03-01	101.98
2022-04-01	105.84
2022-05-01	114.67
2022-06-01	159.57
2022-07-01	123.71

In [9]:

```
1 train_df_oil = train_df['Crude Oil']
2 train_df_oil = pd.DataFrame(train_df_oil)
3 train_df_oil
```

Out[9]:

Crude Oil	
date	
1990-07-01	46.39
1990-08-01	60.68
1990-09-01	87.04
1990-10-01	77.12
1990-11-01	63.04
...	...
2021-01-01	58.31
2021-02-01	68.33
2021-03-01	65.31
2021-04-01	69.62
2021-05-01	72.02

371 rows × 1 columns

In [10]:

```

1 gold = pd.DataFrame()
2
3 fig, ax1 = plt.subplots(figsize=(16,12))
4 ax1.set_title('Gold and Crude Oil Price Over Time', pad = 12)
5 color = 'orange'
6 ax1.set_xlabel('Date')
7 ax1.set_ylabel('Gold', color = color)
8 ax1.plot(train_df.index, train_df['Gold'], color = color)
9 ax1.plot(test_df.index, test_df['Gold'], color = 'red')
10 ax1.tick_params(axis = 'y', labelcolor = color)
11 ax1.yaxis.set_major_formatter(mpl.ticker.StrMethodFormatter('${x:,.0f}'))
12
13 # Adding Twin Axes to plot using dataset_2
14 ax2 = ax1.twinx()
15
16 color = 'black'
17 ax2.set_ylabel('Crude Oil', color = color)
18 ax2.plot(train_df.index, train_df['Crude Oil'], color = color)
19 ax2.plot(test_df.index, test_df['Crude Oil'], color = 'red')
20 ax2.tick_params(axis = 'y', labelcolor = color)
21 ax2.yaxis.set_major_formatter(mpl.ticker.StrMethodFormatter('${x:,.0f}'))

```



In [11]:

```

1 oil_train = train_df['Crude Oil']
2
3
4 oil_train_model = ARIMA(oil_train, order = (0,3,1))
5 oil_train_model_fit = oil_train_model.fit()
6 print(oil_train_model_fit.summary())

```

## SARIMAX Results

```

=====
==
Dep. Variable:          Crude Oil   No. Observations:          3
71
Model:                ARIMA(0, 3, 1)   Log Likelihood            -1332.6
17
Date:                Wed, 13 Jul 2022   AIC                      2669.2
34
Time:                09:15:22         BIC                      2677.0
50
Sample:                07-01-1990      HQIC                     2672.3
40
                        - 05-01-2021
Covariance Type:                opg
=====
==

```

	coef	std err	z	P> z	[0.025	0.97
ma.L1	-0.9998	1.615	-0.619	0.536	-4.165	2.1
sigma2	80.5360	130.201	0.619	0.536	-174.653	335.7

```

-----
--
Ljung-Box (L1) (Q):                59.96   Jarque-Bera (JB):
49.96
Prob(Q):                0.00   Prob(JB):
0.00
Heteroskedasticity (H):                2.62   Skew:
-0.21
Prob(H) (two-sided):                0.00   Kurtosis:
4.75
=====
=====

```

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

C:\Users\wjsdn\AppData\Roaming\Python\Python39\site-packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.  
warn('Non-invertible starting MA parameters found.')

In [12]:

```
1 oil_train_model_fit
```

Out[12]:

<statsmodels.tsa.arima.model.ARIMAResultsWrapper at 0x269afe2a880>

In [13]:

```
1 predict = oil_train_model_fit.predict()
2 oil_train_predict = pd.DataFrame(predict)
3 oil_train_predict
```

Out[13]:

predicted_mean	
date	
1990-07-01	0.000000
1990-08-01	92.781245
1990-09-01	74.383557
1990-10-01	125.463889
1990-11-01	55.097781
...	...
2021-01-01	57.930450
2021-02-01	62.151491
2021-03-01	78.338395
2021-04-01	62.242848
2021-05-01	73.902921

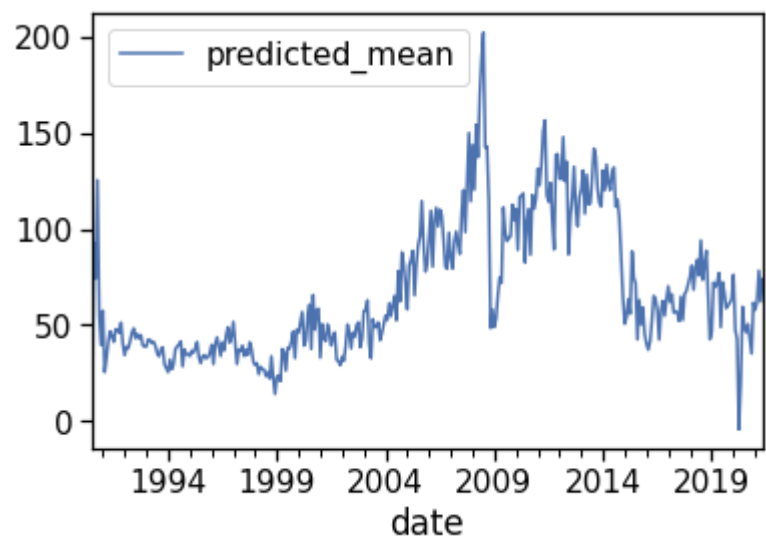
371 rows × 1 columns

In [14]:

```
1 oil_train_predict.plot()
```

Out[14]:

<AxesSubplot:xlabel='date'>



In [15]:

```
1 oil_fore = oil_train_model_fit.forecast(steps=100)
2 oil_fore = pd.DataFrame(oil_fore)
3 oil_fore
```

Out[15]:

predicted_mean	
2021-06-01	74.387812
2021-07-01	76.723436
2021-08-01	79.026871
2021-09-01	81.298119
2021-10-01	83.537178
...	...
2029-05-01	152.552034
2029-06-01	151.829785
2029-07-01	151.075347
2029-08-01	150.288722
2029-09-01	149.469908

100 rows × 1 columns

In [16]:

```
1 oil_fore.idxmax()
```

Out[16]:

```
predicted_mean    2027-07-01  
dtype: datetime64[ns]
```

In [17]:

```
1 oil_fore_error = oil_fore[:14]  
2 oil_fore_error
```

Out[17]:

	<b>predicted_mean</b>
<b>2021-06-01</b>	74.387812
<b>2021-07-01</b>	76.723436
<b>2021-08-01</b>	79.026871
<b>2021-09-01</b>	81.298119
<b>2021-10-01</b>	83.537178
<b>2021-11-01</b>	85.744049
<b>2021-12-01</b>	87.918732
<b>2022-01-01</b>	90.061227
<b>2022-02-01</b>	92.171534
<b>2022-03-01</b>	94.249652
<b>2022-04-01</b>	96.295583
<b>2022-05-01</b>	98.309325
<b>2022-06-01</b>	100.290880
<b>2022-07-01</b>	102.240246



In [18]:

```

1 fig, ax2 = plt.subplots(figsize=(15, 10))
2
3 ax2.set_title('Crude Oil Price Prediction and Actual', pad = 20, fontsize = 30)
4 color = 'black'
5 ax2.set_ylabel('Crude Oil', color = color)
6 ax2.plot(oil_train_predict.index, oil_train_predict['predicted_mean'], color = 'orange')
7 ax2.plot(train_df.index, train_df['Crude Oil'], color = color)
8 ax2.plot(test_df.index, test_df['Crude Oil'], color = 'red')
9 ax2.plot(oil_fore.index, oil_fore['predicted_mean'], color = 'green')
10
11 ax2.tick_params(axis = 'y', labelcolor = color)
12 ax2.yaxis.set_major_formatter(mpl.ticker.StrMethodFormatter('${x:,.0f}'))
13 plt.legend(['Train', 'Actual Price', 'Test', 'Preidcted']);
14 ax2.set_xlabel('date')
15 plt.grid()

```

Crude Oil Price Prediction and Actual



In [19]:

```
1 test_df_oil_error = test_df_oil[1:]  
2 test_df_oil_error.mean()
```

Out[19]:

Crude Oil 96.014286  
dtype: float64

In [20]:

```
1 rms_test = sqrt(mean_squared_error(test_df_oil_error, oil_fore_error))  
2 rms_test
```

Out[20]:

18.645331606904584

In [21]:

```
1 rms_train = sqrt(mean_squared_error(train_df_oil, oil_train_predict))  
2 rms_train
```

Out[21]:

9.645492079222704

In [22]:

```
1 train_df_oil.mean()
```

Out[22]:

Crude Oil 67.93442  
dtype: float64

In [23]:

```
1 from prophet import Prophet
```

In [24]:

```

1 oil_df = all_data['2005-05-01': '2022-07-01']
2 oil_df['ds'] = oil_df.index
3 oil_df.rename(columns = {'Crude Oil' : 'y'}, inplace = True)
4 oil_df = oil_df.reindex(columns=['ds', 'y'])
5 oil_df

```

C:\Users\wjsdn\AppData\Local\Temp\ipykernel\_31800\993763398.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
oil_df['ds'] = oil_df.index
```

C:\Users\wjsdn\AppData\Local\Temp\ipykernel\_31800\993763398.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
oil_df.rename(columns = {'Crude Oil' : 'y'}, inplace = True)
```

Out[24]:

	ds	y
date		
<b>2005-05-01</b>	2005-05-01	78.16
<b>2005-06-01</b>	2005-06-01	84.92
<b>2005-07-01</b>	2005-07-01	90.61
<b>2005-08-01</b>	2005-08-01	102.58
<b>2005-09-01</b>	2005-09-01	97.37
...	...	...
<b>2022-03-01</b>	2022-03-01	101.98
<b>2022-04-01</b>	2022-04-01	105.84
<b>2022-05-01</b>	2022-05-01	114.67
<b>2022-06-01</b>	2022-06-01	159.57
<b>2022-07-01</b>	2022-07-01	123.71

207 rows × 2 columns

In [25]:

```
1 ax = oil_df['y'].plot()  
2 plt.show()
```



In [26]:

```
1 oil_prophet = Prophet(changepoint_prior_scale = 1)  
2 oil_prophet.fit(oil_df)
```

09:15:25 - cmdstanpy - INFO - Chain [1] start processing

09:15:25 - cmdstanpy - INFO - Chain [1] done processing

Out[26]:

<prophet.forecaster.Prophet at 0x269b02ca100>

In [27]:

```
1 forecast_time = 19
2 df_forecast = oil_prophet.make_future_dataframe(periods = forecast_time, freq = 'M')
3 df_forecast
```

Out[27]:

	ds
0	2005-05-01
1	2005-06-01
2	2005-07-01
3	2005-08-01
4	2005-09-01
...	...
221	2023-09-30
222	2023-10-31
223	2023-11-30
224	2023-12-31
225	2024-01-31

226 rows × 1 columns


In [28]:

```
1 df_forecast = oil_prophet.predict(df_forecast)
2 df_forecast
```

Out[28]:


	ds	trend	yhat_lower	yhat_upper	trend_lower	trend_upper	additive_terms	addi
0	2005-05-01	86.586988	72.868965	108.747989	86.586988	86.586988	4.668387	
1	2005-06-01	86.866336	78.650136	115.549001	86.866336	86.866336	9.987831	
2	2005-07-01	87.136672	75.634254	111.273987	87.136672	87.136672	6.853261	
3	2005-08-01	87.416020	74.611630	110.748112	87.416020	87.416020	5.215402	
4	2005-09-01	87.695367	72.139217	109.608403	87.695367	87.695367	3.110068	
...	...	...	...	...	...	...	...	
221	2023-09-30	100.096265	91.004391	149.758995	75.648719	124.544801	19.976464	
222	2023-10-31	100.972340	75.390245	142.297150	73.659790	127.795888	7.105148	
223	2023-11-30	101.820155	66.292015	140.175814	71.800028	132.089511	0.785069	
224	2023-12-31	102.696229	48.309463	123.255013	69.370914	136.154813	-17.041895	
225	2024-01-31	103.572304	47.317726	130.845662	67.072985	140.791943	-15.232677	

226 rows × 16 columns




In [29]:

```
1 oil_prophet.plot(df_forecast, xlabel = 'Year', ylabel = 'Price($)')  
2 #plt.set_major_formatter(mpl.ticker.StrMethodFormatter('${x:,.0f}'))  
3 plt.title('Crude oil Price Prediction by Facebook Prophet', fontsize = 23, pad = 15);
```



In [30]:

```
1 fig, ax = plt.subplots(figsize=(12, 6))
2
3 ax.yaxis.set_major_formatter(mpl.ticker.StrMethodFormatter('${x:,.0f}'))
4
5 plt.plot(df_forecast['ds'], df_forecast['yhat'], label = 'Prediction')
6 plt.plot(df_forecast['ds'], df_forecast['trend_upper'], label = 'Trend Line' )
7 plt.plot(oil_df['ds'], oil_df['y'], label = 'Actual Price', color = 'black')
8 plt.legend()
9 plt.grid()
10 plt.ylabel('Price')
11 plt.xlabel('Year')
12 plt.title('Crude Oil Price Prediction and Trend Line', fontsize = 25, pad = 15);
```






In [31]:

```
1 df_forecast
```

Out[31]:

	ds	trend	yhat_lower	yhat_upper	trend_lower	trend_upper	additive_terms	addi
0	2005-05-01	86.586988	72.868965	108.747989	86.586988	86.586988	4.668387	
1	2005-06-01	86.866336	78.650136	115.549001	86.866336	86.866336	9.987831	
2	2005-07-01	87.136672	75.634254	111.273987	87.136672	87.136672	6.853261	
3	2005-08-01	87.416020	74.611630	110.748112	87.416020	87.416020	5.215402	
4	2005-09-01	87.695367	72.139217	109.608403	87.695367	87.695367	3.110068	
...	...	...	...	...	...	...	...	
221	2023-09-30	100.096265	91.004391	149.758995	75.648719	124.544801	19.976464	
222	2023-10-31	100.972340	75.390245	142.297150	73.659790	127.795888	7.105148	
223	2023-11-30	101.820155	66.292015	140.175814	71.800028	132.089511	0.785069	
224	2023-12-31	102.696229	48.309463	123.255013	69.370914	136.154813	-17.041895	
225	2024-01-31	103.572304	47.317726	130.845662	67.072985	140.791943	-15.232677	

226 rows × 16 columns




In [32]:

```
1 df_forecast_error = df_forecast[:207]
2 df_forecast_error
```

Out[32]:

	ds	trend	yhat_lower	yhat_upper	trend_lower	trend_upper	additive_terms	additi
0	2005-05-01	86.586988	72.868965	108.747989	86.586988	86.586988	4.668387	
1	2005-06-01	86.866336	78.650136	115.549001	86.866336	86.866336	9.987831	
2	2005-07-01	87.136672	75.634254	111.273987	87.136672	87.136672	6.853261	
3	2005-08-01	87.416020	74.611630	110.748112	87.416020	87.416020	5.215402	
4	2005-09-01	87.695367	72.139217	109.608403	87.695367	87.695367	3.110068	
...	...	...	...	...	...	...	...	
202	2022-03-01	83.761709	69.182633	106.381322	83.761709	83.761709	2.407280	
203	2022-04-01	84.637784	71.777500	108.942321	84.637784	84.637784	6.663959	
204	2022-05-01	85.485598	71.326503	108.076995	85.485598	85.485598	4.116755	
205	2022-06-01	86.361673	78.401552	115.016823	86.361673	86.361673	10.024304	
206	2022-07-01	87.209488	75.653664	112.018439	87.209488	87.209488	7.085945	

207 rows × 16 columns



In [33]:

```
1 facebook_rmse = sqrt(mean_squared_error(oil_df['y'], df_forecast_error['yhat']))
2 facebook_rmse
```

Out[33]:

14.08660156768478

In [ ]:

```
1
```